

WES 237A: Introduction to Embedded System Design (Winter 2024)

Lab 3: Serial and CPU

Due: 2/4/2024 11:59pm

Ricardo Lizarraga, rlizarraga0@gmail.com
619.252.4157

PID: A69028483


<https://github.com/RiLizarraga>

In order to report and reflect on your WES 237A labs, please complete this Post-Lab report by the end of the weekend by submitting the following 2 parts:

- Upload your lab 3 report composed by a single PDF that includes your in-lab answers to the bolded questions in the Google Doc Lab and your Jupyter Notebook code. You could either scan your written copy, or simply type your answer in this Google Doc. **However, please make sure your responses are readable.**
 - Answer two short essay-like questions on your Lab experience.
- All responses should be submitted to Canvas. Please also be sure to push your code to your git repo as well.

Serial Connection

- Using a micro USB cable, connect your board to your laptop
- Connect to board using the serial connection
 - Linux
 - Open a new terminal
 - Run the command
 - `sudo screen /dev/<port> 115200 #port: ttyUSB0 or ttyUSB1`
 - MAC
 - Open a new terminal
 - Run the command and check the PYNQ resources for the port
 - `sudo screen /dev/<port> 115200 #port: check resources`
 - Windows
 - Check the resource for how to connect through serial to the PYNQ board
 - Resources:
 - https://pynq.readthedocs.io/en/v2.0/getting_started.html
 - <https://www.nengo.ai/nengo-pynq/connect.html>
 - After connecting



Boot loader terminal commands

zynq> help		icache	- enable or disable instruction cache
?	- alias for 'help'	iminfo	- print header information for application image
base	- print or set address offset	imls	- list all images found in flash
bdinfo	- print Board Info structure	imxtract	- extract a part of a multi-image
blkcache	- block cache diagnostics and control	itest	- return true/false on integer compare
boot	- boot default, i.e., run 'bootcmd'	led	- manage LEDs
bootd	- boot default, i.e., run 'bootcmd'	ln	- Create a symbolic link
bootefi	- Boots an EFI payload from memory	load	- load binary file from a filesystem
bootelf	- Boot from an ELF image in memory	loadb	- load binary file over serial line (kermit mode)
bootm	- boot application image from memory	loads	- load S-Record file over serial line
bootp	- boot image via network using BOOTP/TFTP protocol	loadx	- load binary file over serial line (xmodem mode)
bootvx	- Boot vxWorks from an ELF image	loady	- load binary file over serial line (ymodem mode)
bootz	- boot Linux zImage image from memory	loop	- infinite loop on address range
chpart	- change active partition of a MTD device	ls	- list files in a directory (default /)
clk	- CLK sub-system	md	- memory display
cmp	- memory compare	mdio	- MDIO utility commands
coninfo	- print console devices and information	mii	- MII utility commands
cp	- memory copy	mm	- memory modify (auto-incrementing address)
crc32	- checksum calculation	mmc	- MMC sub system
dcache	- enable or disable data cache	mmcinfo	- display MMC info
dfu	- Device Firmware Upgrade	mtd	- MTD utils
dhcp	- boot image via network using DHCP/TFTP protocol	mtddparts	- define flash/nand partitions
dm	- Driver model low level access	mtest	- simple RAM read/write test
echo	- echo args to console	mw	- memory write (fill)
editenv	- edit environment variable	nand	- NAND sub-system
efidebug	- Configure UEFI environment	nboot	- boot from NAND device
env	- environment handling commands	net	- NET sub-system
erase	- erase FLASH memory	nfs	- boot image via network using NFS protocol
exit	- exit script	nm	- memory modify (constant address)
ext2load	- load binary file from a Ext2 filesystem	panic	- Panic with optional message
ext2ls	- list files in a directory (default /)	part	- disk partition related commands
ext4load	- load binary file from a Ext4 filesystem	ping	- send ICMP ECHO_REQUEST to network host
ext4ls	- list files in a directory (default /)	printenv	- print environment variables
ext4size	- determine a file's size	protect	- enable or disable FLASH write protection
ext4write	- create a file in the root directory	pxe	- commands to get and boot from pxe files
false	- do nothing, unsuccessfully	random	- fill memory with random pattern
fatinfo	- print information about filesystem	reset	- Perform RESET of the CPU
fatload	- load binary file from a dos filesystem	run	- run commands in an environment variable
fatls	- list files in a directory (default /)	save	- save file to a filesystem
fatmkdir	- create a directory	saveenv	- save environment variables to persistent storage
fatrm	- delete a file	setenv	- set environment variables
fatsize	- determine a file's size	sf	- SPI flash sub-system
fatwrite	- write file into a dos filesystem	showvar	- print local hushshell variables
fdt	- flattened device tree utility commands	size	- determine a file's size
flinfo	- print FLASH memory information	sleep	- delay execution for some time
fpga	- loadable FPGA image support	source	- run script from memory
fru	- FRU table info	spl	- SPL configuration
fstype	- Look up a filesystem type	sysboot	- command to get and boot from syslinux files
fstypes	- List supported filesystem types	test	- minimal test like /bin/sh
go	- start application at address 'addr'	tftpboot	- boot image via network using TFTP protocol
gpio	- query and control gpio pins	tftpboot	- TFTP put command, for uploading files to a server
help	- print command description/usage	thor	- TIZEN "THOR" downloader
i2c	- I2C sub-system	time	- run commands and summarize execution time
		timer	- access the system timer
		true	- do nothing, successfully
		ubi	- ubi commands
		ubifsload	- load file from an UBIFS filesystem
		ubifsls	- list files in a directory
		ubifsmount	- mount UBIFS volume
		ubifsunmount	- unmount UBIFS volume
		usb	- USB sub-system
		usbboot	- boot from USB device
		version	- print monitor, compiler and linker version
		zynq	- Zynq specific commands

- o Restart the board (\$ *sudo reboot*)
- o Interrupt the boot (*keyboard interrupt*)
- o List current settings (*printenv*)
- o Put a screenshot of your \$ *printenv* output


```

Zynq> printenv
arch=arm
baudrate=115200
board=zynq
board_name=zynq
boot_a_script=load ${devtype} ${devnum}:${distro_bootpart} ${scriptaddr} ${prefix}${script}; source ${scriptaddr}
boot_efi_binary=load ${devtype} ${devnum}:${distro_bootpart} ${kernel_addr_r} efi/boot/bootarm.efi; if fdt addr ${fdt_addr_r}; th
en bootefi ${kernel_addr_r} ${fdt_addr_r}; else bootefi ${kernel_addr_r} ${fdtcontroladdr}; fi
boot_efi_bootmgr=if fdt addr ${fdt_addr_r}; then bootefi bootmgr ${fdt_addr_r}; else bootefi bootmgr; fi
boot_extlinux=sysboot ${devtype} ${devnum}:${distro_bootpart} any ${scriptaddr} ${prefix}${boot_syslinux_conf}
boot_net usb start=usb start
boot_prefixes=/ /boot/
boot_script_dhcp=boot.scr.uimg
boot_scripts=boot.scr.uimg boot.scr
boot_syslinux_conf=extlinux/extlinux.conf
boot_targets=mmc0 jtag mmc0 mmcl qspi nand nor usb0 usb1 pxe dhcp
bootcmd=run distro bootcmd
bootcmd_dhcp=devtype=dhcp; run boot_net usb start; if dhcp ${scriptaddr} ${boot_script_dhcp}; then source ${scriptaddr}; fi; seten
v efi_fdtfile ${fdtfile}; if test -z "${fdtfile}" -a -n "${soc}"; then setenv efi_fdtfile ${soc}-${board}${boardver}.dtb; fi; set
env efi_old_vci ${bootp_vci}; setenv efi_old_arch ${bootp_arch}; setenv bootp_vci PXEClient:Arch:00010:UNDI:003000; setenv bootp_ar
ch 0xai; if dhcp ${kernel_addr_r}; then tftpboot ${fdt_addr_r} dtb/${efi_fdtfile}; if fdt addr ${fdt_addr_r}; then bootefi ${kernel_a
ddr_r} ${fdt_addr_r}; else bootefi ${kernel_addr_r} ${fdtcontroladdr}; fi; fi; setenv bootp_vci ${efi_old_vci}; setenv bootp_arch ${e
fi_old_arch}; setenv efi_fdtfile; setenv efi_old_arch; setenv efi_old_vci;
bootcmd_jtag=echo JTAG: Trying to boot script at ${scriptaddr} && source ${scriptaddr}; echo JTAG: SCRIPT FAILED: continuing...;
bootcmd_mmc0=devnum=0; run mmc boot
bootcmd_mmcl=devnum=1; run mmc boot
bootcmd_nand=nand info && nand read ${scriptaddr} ${script_offset_f} ${script_size_f} && echo NAND: Trying to boot script at ${sc
riptaddr} && source ${scriptaddr}; echo NAND: SCRIPT FAILED: continuing...;
bootcmd_nor=cp.b ${script_offset_nor} ${scriptaddr} ${script_size_f} && echo NOR: Trying to boot script at ${scriptaddr} && sourc
e ${scriptaddr}; echo NOR: SCRIPT FAILED: continuing...;
bootcmd_pxe=run boot_net usb start; dhcp; if pxe get; then pxe boot; fi
bootcmd_qspi=sf probe 0 0 0 && sf read ${scriptaddr} ${script_offset_f} ${script_size_f} && echo QSPI: Trying to boot script at $
{scriptaddr} && source ${scriptaddr}; echo QSPI: SCRIPT FAILED: continuing...;
bootcmd_usb0=devnum=0; run usb boot
bootcmd_usb1=devnum=1; run usb boot
bootcmd_usb_dfu0=setenv dfu_alt_info boot.scr ram ${scriptaddr} ${script_size_f} && dfu 0 ram 0 60 && echo DFU0: Trying to boot scrip
t at ${scriptaddr} && source ${scriptaddr}; echo DFU0: SCRIPT FAILED: continuing...;
bootcmd_usb_dfu1=setenv dfu_alt_info boot.scr ram ${scriptaddr} ${script_size_f} && dfu 1 ram 1 60 && echo DFU1: Trying to boot scrip
t at ${scriptaddr} && source ${scriptaddr}; echo DFU1: SCRIPT FAILED: continuing...;
bootcmd_usb_thor0=setenv dfu_alt_info boot.scr ram ${scriptaddr} ${script_size_f} && thordown 0 ram 0 && echo THOR0: Trying to boot s
cript at ${scriptaddr} && source ${scriptaddr}; echo THOR0: SCRIPT FAILED: continuing...;
bootcmd_usb_thor1=setenv dfu_alt_info boot.scr ram ${scriptaddr} ${script_size_f} && thordown 1 ram 1 && echo THOR1: Trying to boot s
cript at ${scriptaddr} && source ${scriptaddr}; echo THOR1: SCRIPT FAILED: continuing...;
bootdelay=2
bootm_low=0
bootm_size=20000000
cpu=armv7
distro_bootcmd=for target in ${boot_targets}; do run bootcmd_${target}; done
efi_dtb_prefixes=/ /dtb/ /dtb/current/
ethaddr=00:00:05:6b:03:3a
fdt_addr_r=0x1f00000
fdtcontroladdr=1eae1160
kernel_addr_r=0x2000000
load_efi_dtb=load ${devtype} ${devnum}:${distro_bootpart} ${fdt_addr_r} ${prefix}${efi_fdtfile}
loadaddr=0x0
mmc boot=if mmc dev ${devnum}; then devtype=mmc; run scan_dev_for_boot_part; fi
modeboot=sdboot
pxefile_addr_r=0x2000000
ramdisk_addr_r=0x3100000
scan_dev_for_boot=echo Scanning ${devtype} ${devnum}:${distro_bootpart}...; for prefix in ${boot_prefixes}; do run scan_dev_for_e
xtlinux; run scan_dev_for_scripts; done; run scan_dev_for_efi;
scan_dev_for_boot_part=part list ${devtype} ${devnum} -bootable devplist; env exists devplist || setenv devplist 1; for distro_bo
otpart in ${devplist}; do if fstype ${devtype} ${devnum}:${distro_bootpart} bootfstype; then run scan_dev_for_boot; fi; done; set
env devplist
scan_dev_for_efi=setenv efi_fdtfile ${fdtfile}; if test -z "${fdtfile}" -a -n "${soc}"; then setenv efi_fdtfile ${soc}-${board}${
boardver}.dtb; fi; for prefix in ${efi_dtb_prefixes}; do if test -e ${devtype} ${devnum}:${distro_bootpart} ${prefix}${efi_fdtfil
e}; then run load_efi_dtb; fi; done; run boot_efi_bootmgr; if test -e ${devtype} ${devnum}:${distro_bootpart} efi/boot/bootarm.efi;
then echo Found EFI removable media binary efi/boot/bootarm.efi; run boot_efi_binary; echo EFI LOAD FAILED: continuing...; fi; se
tenv efi_fdtfile
scan_dev_for_extlinux=if test -e ${devtype} ${devnum}:${distro_bootpart} ${prefix}${boot_syslinux_conf}; then echo Found ${prefix
}${boot_syslinux_conf}; run boot_extlinux; echo SCRIPT FAILED: continuing...; fi
scan_dev_for_scripts=for script in ${boot_scripts}; do if test -e ${devtype} ${devnum}:${distro_bootpart} ${prefix}${script}; the
n echo Found U-Boot script ${prefix}${script}; run boot_a_script; echo SCRIPT FAILED: continuing...; fi; done
script_offset_f=fc0000
script_offset_nor=0xe2fc0000
script_size_f=0x40000
scriptaddr=3000000
soc=zynq
stderr=serial@e0000000
stdin=serial@e0000000
stdout=serial@e0000000
ubifs_boot=env exists bootubipart || env set bootubipart UBI; env exists bootubivol || env set bootubivol boot; if ubi part ${boo
tubipart} && ubismount ubi${devnum}:${bootubivol}; then devtype=ubi; run scan_dev_for_boot; fi
usb boot=usb start; if usb dev ${devnum}; then devtype=usb; run scan_dev_for_boot_part; fi
vendor=xilinx

```

Environment size: 5849/131067 bytes

Zynq> █

```

0[#####*****          58.0%]   Tasks: 27, 8 thr; 2 running
1[#####*****          42.9%]   Load average: 3.54 1.19 0.42
Mem[|||||##*****98.9M/494M]   Uptime: 00:01:29
Swp[|                    256K/512M]

PID USER      PRI  NI  VIRT   RES   SHR  S  CPU%  VMEM%   TIME+  Command
879 root        20   0 31828  5084  2972 R  10.6  1.0   0:00.16 (haveged)
  1 root        20   0 31824  7804  5696 S   7.9  1.5   0:07.37 /sbin/init earlylp
110 root        19  -1 51928 13072 12308 S   1.3  2.6   0:01.44 /lib/systemd/syst
651 xilinx      20   0  7844  2688  2160 R   1.3  0.5   0:00.89 htop
262 messagebu  20   0  6684  3020  2520 S   0.7  0.6   0:00.60 @dbus-daemon --sy
295 syslog     20   0 25768  2940    2
159 root        20   0 18680  3724  3012 S   0.0  0.7   0:00.67 /lib/systemd/syst
173 systemd-r  20   0 18872  8704  5628 S   0.0  1.7   0:00.83 /lib/systemd/syst
174 systemd-t  20   0 19960  2732  2232 S   0.0  0.5   0:00.67 /lib/systemd/syst
208 systemd-t  20   0 19960  2732  2232 S   0.0  0.5   0:00.05 /lib/systemd/syst
237 root        20   0 32520  4440  3516 S   0.0  0.9   0:00.05 /sbin/dhclient -l
244 root        20   0 32520  4440  3516 S   0.0  0.9   0:00.01 /sbin/dhclient -l
245 root        20   0 32520  4440  3516 S   0.0  0.9   0:00.00 /sbin/dhclient -l
246 root        20   0 32520  4440  3516 S   0.0  0.9   0:00.00 /sbin/dhclient -l
261 root        20   0  5184   924   820 S   0.0  0.2   0:00.02 /usr/sbin/cron -f

F1Help  F2Setup  F3Search  F4Filter  F5Tree  F6SortBy  F7Nice  -F8Nice  +F9Kill  F10Quit

```


Change Bootargs

- If you need to return to the default bootargs, you can find them below
 - https://github.com/Xilinx/PYNQ/blob/master/sdbuild/boot/meta-pynq/recipes-bsp/device-tree/files/pynq_bootargs.dtsi
 - `bootargs = 'root=/dev/mmcblk0p2 rw earlyprintk rootfstype=ext4 rootwait devtmpfs.mount=1 uio_pdrv_genirq.of_id="generic-uio" clk_ignore_unused'`
- To edit bootargs:
 - Interrupt the boot {during initia sl boot, just press any key to enter into boot loader}
 - Edit boot arguments:
 - `$ editenv bootargs` {before making changes to bootargs variable we can see the actual value by `$ echo $bootargs`}
 - Insert arguments included the quotations all in one line:
 - Bootargs (default and more) are at [here](#)

```
bootargs = 'console=ttyPS0,115200 root=/dev/mmcblk0p2 rw earlyprintk rootfstype=ext4 rootwait
devtmpfs.mount=1 uio_pdrv_genirq.of_id="generic-uio" clk_ignore_unused isolcpus=1 && bootz
0x03000000 - 0x02A00000'
```

- `$ boot`
 - Change bootargs to the following
 - `bootargs = 'console=ttyPS0,115200 root=/dev/mmcblk0p2 rw earlyprintk rootfstype=ext4 rootwait devtmpfs.mount=1 uio_pdrv_genirq.of_id="generic-uio" clk_ignore_unused isolcpus=1 && bootz 0x03000000 - 0x02A00000'`

```
Zynq> edit bootargs
edit: bootargs = 'console=ttyPS0,115200 root=/dev/mmcblk0p2 rw earlyprintk rootfstype=ext4 rootwait
devtmpfs.mount=1 uio_pdrv_genirq.of_id="generic-uio" clk_ignore_unused isolcpus=1 && bootz 0x03000000 - 0x02A00000'
Zynq> echo $bootargs
bootargs = console=ttyPS0,115200 root=/dev/mmcblk0p2 rw earlyprintk rootfstype=ext4 rootwait devtmpfs.mount=1 uio_pdrv_genirq.of_id="generic-uio" clk_ignore_unused isolcpus=1 && bootz 0x03000000 - 0x02A00000
```

- What does `isolcpus=1` do?

Remove the specified CPUs, as defined by the `cpu_number` values. These are set of CPUs that the kernel process scheduler will not interact.

See [htop](#) displays No utilization of CPU1, while everything executes on CPU0:

COM5 - PuTTY

```
0[#####100.0%] Tasks: 28, 10 thr; 2 running
1[0.0%] Load average: 0.42 0.25 0.32
Mem[|||||]##*****124M/494M Uptime: 00:14:47
Swp[0K/512M]

  PID USER      PRI  NI  VIRT   RES   SHR  S  CPU%  MEM%   TIME+  Command
 2040 root        20   0 33940 23808  9752 R  83.8   4.7   0:02.46 /usr/local/share/
   1 root        20   0 30800  7800  5696 S   5.8   1.5   0:12.76 /sbin/init bootar
  575 root        20   0  125M 67160 16864 S   0.6  13.3   0:28.95 /usr/local/share/
 1953 xilinx      20   0  7736  2428  2060 R   0.6   0.5   0:00.78 htop
 2042 root        20   0  125M 67160 16864 S   0.6  13.3   0:00.02 /usr/local/share/
  108 root        19  -1 60208 13432 12636 S   0.0   2.7   0:02.15 /lib/systemd/syst
  147 root        20   0 18552  3672  3020 S   0.0   0.7   0:00.59 /lib/systemd/syst
  171 systemd-r    20   0 18872  8708  5628 S   0.0   1.7   0:01.58 /lib/systemd/syst
  173 systemd-t    20   0 19960  2732  2232 S   0.0   0.5   0:01.17 /lib/systemd/syst
```

- What would `isolcpus=0` do?

The kernel process scheduler will interact with all CPU's (in this case CPU0 & CPU1)-shared Load

Heavy CPU Utilization

- Download *fib.py* from [here](#). This is a recursive implementation for generating Fibonacci sequences. We just do not print the results.
 - Make sure your board is booted with custom bootargs above including *isolcpus=1*
- 1) Open two terminals (Jupyter):
- Terminal 1: run *htop* to monitor CPU utilization

```
0[||||| 6.5%] Tasks: 29, 21 thr; 1 running
1[ 0.0%] Load average: 0.00 0.00 0.00
Mem[||||| 142M/494M] Uptime: 01:00:45
Swp[ 0K/512M]
```

- Terminal 2: run *\$ python3 fib.py* and monitor CPU utilization and time spent for running the script (set terms to lower than 40)

```
root@pynq:/home/xilinx/jupyter_notebooks/RLS/Lab3# python3 fib.py
How many terms? 40
```

- Describe the results of *htop*.

```
0[||||| 100.0%] Tasks: 30, 21 thr; 2 running
1[ 0.0%] Load average: 0.04 0.05 0.00
Mem[||||| 143M/494M] Uptime: 01:05:51
Swp[ 0K/512M]
```

The CPU0's usage jumped to a 100% usage because the intensive processing of *fib.py*

PID	USER	PRI	NI	VIRT	RES	SHR	S	CPU%	MEM%	TIME+	Command
6421	root	20	0	13052	6036	4120	R	94.2	1.2	2:24.50	python3 fib.py

{noticed the CPU1 remains 0% (no use at all)}

- 2) Repeat the previous part, but this time use *taskset* to use CPU1:
- Terminal 2: run *\$ taskset -c 1 python3 fib.py* and monitor CPU utilization and time spent for running the script

```
root@pynq:/home/xilinx/jupyter_notebooks/RLS/Lab3# taskset -c 1 python3 fib.py
How many terms? 40
```

- Describe the results of *htop*. Specifically, what's different from running it in 1)?

Since this time we attached the process "python3 fib.py" to CPU1 by issuing the command:

\$ taskset -c 1 python3 fib.py

See the heavy use of CPU1 which is running *python3 fib.py*, WHILE CPU0 is lighter usage

```
0[||| 2.0%] Tasks: 31, 21 thr; 2 running
1[||||| 100.0%] Load average: 0.71 0.64 0.33
Mem[||||| 146M/494M] Uptime: 01:12:34
Swp[ 0K/512M]
```

3) Heavy Utilization on CPU0:

- Open another terminal and run *\$ dd if=/dev/zero of=/dev/null*

```
root@pynq:/# dd if=/dev/zero of=/dev/null
```

- Repeat parts 1 and 2
- Describe the results of *htop*.

```
0[||||| 100.0%] Tasks: 33, 21 thr; 2 running
1[||||| 100.0%] Load average: 1.39 1.01 0.61
Mem[||||| 147M/494M] Uptime: 01:19:06
Swp[ 0K/512M]
```

After issuing command "***\$ dd if=/dev/zero of=/dev/null***", now the usage load is shared between CPU0 & CPU1

ARM Performance Monitoring (C++)

- Download [kernel_modules](#) folder
- Read through CPUcntr.c and reference the ARM documentation for the PMU registers [here](#) to answer the following question.
 - According to the ARM docs, what does the following line do? Are they written in assembly code, python, C, or C++?

```
■ asm("MCR p15, 0, 1, c9, c14, 0\n\t");
```

This code is Assembly language;

```
asm("MCR p15, 0, 1, c9, c14, 0\n\t"); {program the performance-counter control-register}
```

The PMU module enables users to access performance counter, Providing Access to PMU for Applications on both CPUs, this kernel object can enable user-mode for PMU.

ARM Performance Monitoring Unit (PMU)

- PMU is an event counting hardware which can be used to profile and benchmark code
- Cortex-A9 PMU provides six counters
- Each counter can count any of the 58 events available in the Cortex-A9 processor
- These counters can be accessed through debugging tools or directly through the CP15 Registers

Accessing ARM PMU

- By default user space (applications), do not have access to PMU
- Kernel has access to PMU

MCR: Move to Coprocessor from ARM Register, PMU coprocID is p15

System control processor registers overview

This section gives details of all the registers in the system control coprocessor. The section presents a summary of the registers and detailed descriptions in register order of CRn, Opcode_1, CRm, Opcode_2.

You can access CP15 registers with MRC and MCR instructions, as described in [Use of the system control coprocessor](#):

```
MCR{cond} p15, <Opcode_1>, <Rd>, <CRn>, <CRm>, <Opcode_2>
```

```
MRC{cond} p15, <Opcode_1>, <Rd>, <CRn>, <CRm>, <Opcode_2>
```

You can access the ARM PMU (ARM Performance Monitor Unit) counters, and their associated control registers through the internal CP15 interface, and through the APB, using the relevant offset when PADDRDBG[12]=1

- Compile and insert the kernel module following the instructions from the README file.
- Download [clock_example](#) folder
- Read through include/cycletime.h and take note of the functions to initialize the counters and get the cyclecount (what datatype do they return, what parameters do they take)
- What does the following line do?

```
■ asm volatile ("MRC p15, 0, %0, c9, c13, 0\n\t" : "=r"(value));
```

This line will return the cyclecount

```
{{{ Per Teacher Assistant, we stop here because compiler change issues with the PYNQ board}}}
```

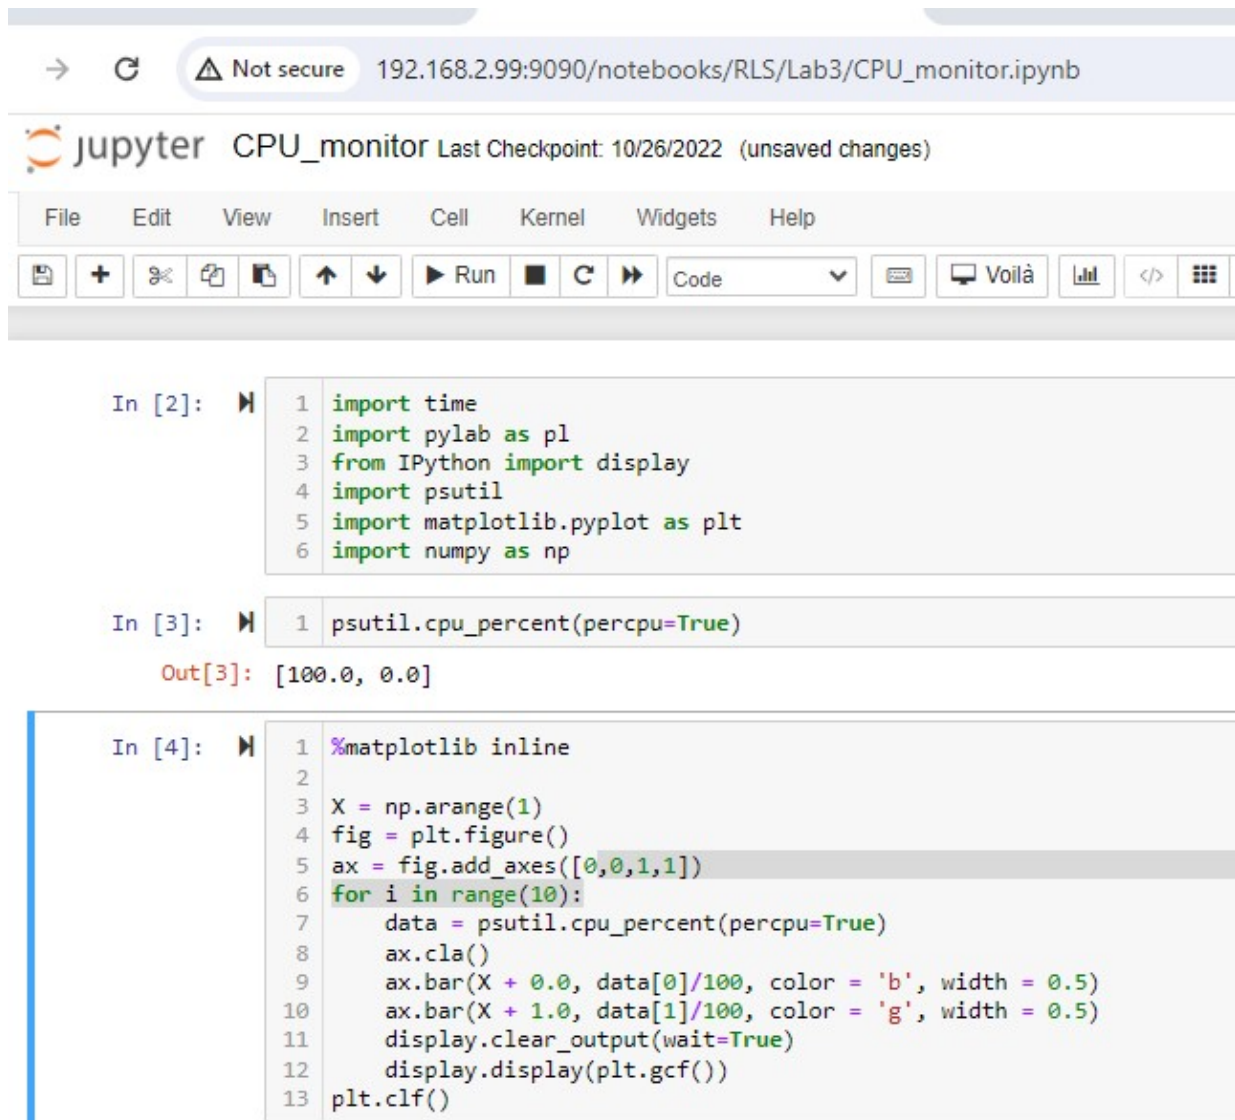
Complete the code in src/main.cpp. These instructions are for those who have never coded in C++

- Declare 2 variables (cpu_before, cpu_after) of the correct datatype
- Initialize the counter
- Get the cyclecount 'before' sleeping
- Get the cyclecount 'after' sleeping
- Print the difference number of counts between starting and stopping the counter
- After completing the code, open a jupyter terminal and change directory to clock_examples/
- Run `$ make` to compile the code
- Run the code with `$./lab4 <delay-time-seconds>`
- Change the delay time and note down the different cpu cycles as well as the different timers.

```
{{{ Per Teacher Assistant, we stop here because compiler change issues with the PYNQ board}}}
```

Jupyter Notebook CPU Monitoring (OPTIONAL)

- Download CPU_monitor.ipynb from [here](#). This is an interactive implementation for plotting in a loop. Running this notebook is a computationally heavy task for your CPU, therefore you do not need to run any additional process to utilize your CPU0.



The screenshot shows a Jupyter Notebook interface in a web browser. The address bar indicates a local connection to 192.168.2.99:9090. The notebook is titled 'CPU_monitor' and shows the last checkpoint from 10/26/2022. The interface includes a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for file operations, running cells, and viewing output. The code is organized into three input cells:

```
In [2]: 1 import time
        2 import pylab as pl
        3 from IPython import display
        4 import psutil
        5 import matplotlib.pyplot as plt
        6 import numpy as np

In [3]: 1 psutil.cpu_percent(percpu=True)

Out[3]: [100.0, 0.0]


In [4]: 1 %matplotlib inline
        2
        3 X = np.arange(1)
        4 fig = plt.figure()
        5 ax = fig.add_axes([0,0,1,1])
        6 for i in range(10):
        7     data = psutil.cpu_percent(percpu=True)
        8     ax.cla()
        9     ax.bar(X + 0.0, data[0]/100, color = 'b', width = 0.5)
       10     ax.bar(X + 1.0, data[1]/100, color = 'g', width = 0.5)
       11     display.clear_output(wait=True)
       12     display.display(plt.gcf())
       13 plt.clf()
```

- Create a Jupyter notebook
 - Use the `os` library to create a python program that accepts a number from user input (0 or 1) and runs `fib.py` on a specific core (0 or 1).
 - Hint: Look at the `os.system()` call and remember the 'taskset' function we've used previously.

```
1 import os
2 cpu = input('enter CPU number (0/1)?')
3 str = "taskset -c "+cpu+" python3 fib.py"
4 print(str)
5 os.system(str)
```


- You should have two notebooks running: 1) CPU_monitor, 2) CPU_select
- Compare your observations between using Jupyter notebook CPU_monitor and linux command *htop* for monitoring CPU utilization.

Before the test:




```
0[|||||||||||||||||||||||||||||||||99.4%] Tasks: 29, 32 thr; 2 running
1[||| 6.5%] Load average: 0.58 0.22 0.19
```

CPU1 running python3 fib.py:

```
[*]: 1 import os
      2 cpu = input('enter CPU number (0/1)?')
      3 str = "taskset -c "+cpu+" python3 fib.py"
      4 print(str)
      5 os.system(str)
```

```
enter CPU number (0/1)?1
taskset -c 1 python3 fib.py
```



```
|||||||||||||||||||||||||||||||||99.4% 31 2 2
1[|||||||||||||||||||||||||||||100.0%] Load average: 1.17 0.69 0.40
```

HTOP application on console responses significantly faster than the monitoring application on Jupyter notebook, possible reasons are that htop app maybe a c++ compiled application, while the python jupyter code is translated code into c++ libraries on the fly, all these plus the graphical overhead takes a lot of CPU/ Clock resources